

Decoupled Neural Interfaces for PyTorch

This tiny library is an implementation of Decoupled Neural Interfaces using Synthetic Gradients

<<https://arxiv.org/abs/1608.05343>> _ for PyTorch <<http://pytorch.org/>> _ . It's very simple to use as it was designed to enable researchers to integrate DNI into existing models with minimal amounts of code.

To install, run::

```
$ python setup.py install
```

Description of the library and how to use it in some typical cases is provided below. For more information, please read the code.

Terminology

This library uses a message passing abstraction introduced in the paper. Some terms used in the API (matching those used in the paper wherever possible):

- `Interface` - A Decoupled Neural Interface that decouples two parts (let's call them part A and part B) of the network and lets them communicate via `message` passing. It may be `Forward` , `Backward` Or `Bidirectional` .
- `BackwardInterface` - A type of `Interface` that the paper focuses on. It can be used to prevent update locking by predicting gradient for part A of the decoupled network based on the activation of its last layer.
- `ForwardInterface` - A type of `Interface` that can be used to prevent forward locking by predicting input for part B of the network based on some information known to both parts - in the paper it's the input of the whole network.
- `BidirectionalInterface` - A combination of `ForwardInterface` and `BackwardInterface` , that can be used to achieve a complete unlock.
- `message` - Information that is passed through an `Interface` - activation of the last layer for `ForwardInterface` or gradient w.r.t. that activation for `BackwardInterface` . Note that no original information passes through. A `message` is consumed by one end of the `Interface` and used to update a `Synthesizer` . Then the `Synthesizer` can be used produce a synthetic `message` at the other end of the `Interface` .
- `trigger` - Information based on which `message` is synthesized. It needs to be accessible by both parts of the network. For `BackwardInterface` , it's activation of the layer w.r.t. which gradient is to be synthesized. For `ForwardInterface` it can be anything - in the paper it's the input of the whole network.
- `context` - Additional information normally not shown to the network at the forward pass, that can condition an `Interface` to provide a better estimate of the `message` . The paper uses labels for this purpose and calls DNI with context `cDNI`.
- `send` - A method of an `Interface` , that takes as input `message` and `trigger` , based on which that `message` should be generated, and updates `Synthesizer` to improve the estimate.
- `receive` - A method of an `Interface` , that takes as input `trigger` and returns a `message` generated by a `Synthesizer` .
- `Synthesizer` - A regression model that estimates `message` based on `trigger` and `context` .

Typical use cases

Synthetic Gradient for Feed-Forward Networks ^^^

In this case we want to decouple two parts A and B of a neural network to achieve an update unlock, so that there is a normal forward pass from part A to B, but part A learns using synthetic gradient generated by the DNI.

.. image:: images/feedforward-update-unlock.png

Following the paper's convention, solid black arrows are update-locked forward connections, dashed black arrows are update-unlocked forward connections, green arrows are real error gradients and blue arrows are synthetic error gradients. Full circles denote synthetic gradient loss computation and `Synthesizer` update.

We can use a `BackwardInterface` to do that:

.. code-block:: python

```
class Network(torch.nn.Module):

    def __init__(self):
        # ...

        # 1. create a BackwardInterface, assuming that dimensionality of
        # the activation for which we want to synthesize gradients is
        # activation_dim
        self.backward_interface = dni.BackwardInterface(
            dni.BasicSynthesizer(output_dim=activation_dim, n_hidden=1)
        )

        # ...

    def forward(self, x):
        # ...

        # 2. call the BackwardInterface at the point where we want to
        # decouple the network
        x = self.backward_interface(x)

        # ...

        return x
```

That's it! During the forward pass, `BackwardInterface` will use a `Synthesizer` to generate synthetic gradient w.r.t. activation, backpropagate it and add to the computation graph a node that will intercept the real gradient during the backward pass and use it to update the `Synthesizer`'s estimate.

The `Synthesizer` used here is `BasicSynthesizer` - a multi-layer perceptron with ReLU activation function. Writing a custom `Synthesizer` is described at [Writing custom Synthesizers](#).

You can specify a `context` by passing `context_dim` (dimensionality of the context vector) to the `BasicSynthesizer` constructor and wrapping all DNI calls in the `dni.synthesizer_context` context manager:

.. code-block:: python

```
class Network(torch.nn.Module):

    def __init__(self):
        # ...

        self.backward_interface = dni.BackwardInterface(
            dni.BasicSynthesizer(
                output_dim=activation_dim, n_hidden=1,
                context_dim=context_dim
            )
        )

        # ...

    def forward(self, x, y):
        # ...
```

```

# assuming that context is labels given in variable y
with dni.synthesizer_context(y):
    x = self.backward_interface(x)

# ...

return x

```

Example code for digit classification on MNIST is at `examples/mnist-mlp` <`examples/mnist-mlp`> _.

Complete Unlock for Feed-Forward Networks ^^^

In this case we want to decouple two parts A and B of a neural network to achieve forward and update unlock, so that part B receives synthetic input and part A learns using synthetic gradient generated by the DNI.

.. image:: images/feedforward-complete-unlock.png

Red arrows are synthetic inputs.

We can use a `BidirectionalInterface` to do that:

.. code-block:: python

```

class Network(torch.nn.Module):

    def __init__(self):
        # ...

        # 1. create a BidirectionalInterface, assuming that dimensionality of
        # the activation for which we want to synthesize gradients is
        # activation_dim and dimensionality of the input of the whole
        # network is input_dim
        self.bidirectional_interface = dni.BidirectionalInterface(
            # Synthesizer generating synthetic inputs for part B, trigger
            # here is the input of the network
            dni.BasicSynthesizer(
                output_dim=activation_dim, n_hidden=1,
                trigger_dim=input_dim
            ),
            # Synthesizer generating synthetic gradients for part A,
            # trigger here is the last activation of part A (no need to
            # specify dimensionality)
            dni.BasicSynthesizer(
                output_dim=activation_dim, n_hidden=1
            )
        )

        # ...

    def forward(self, input):
        x = input

        # ...

        # 2. call the BidirectionalInterface at the point where we want to
        # decouple the network, need to pass both the last activation
        # and the trigger, which in this case is the input of the whole
        # network
        x = self.backward_interface(x, input)

        # ...

        return x

```

During the forward pass, `BidirectionalInterface` will receive real activation, use it to update the input `Synthesizer`, generate synthetic gradient w.r.t. that activation using the gradient `Synthesizer`, backpropagate it, generate synthetic input using the input `Synthesizer` and attach to it a computation graph node that will intercept the real gradient w.r.t. the synthetic input and use it to update the gradient `Synthesizer`.

Example code for digit classification on MNIST is at `examples/mnist-full-unlock` `<examples/mnist-full-unlock>`.

Writing custom Synthesizers ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^

This library includes only `BasicSynthesizer` - a very simple `Synthesizer` based on a multi-layer perceptron with ReLU activation function. It may not be sufficient for all cases, for example for classifying MNIST digits using a CNN the paper uses a `Synthesizer` that is also a CNN.

You can easily write a custom `Synthesizer` by subclassing `torch.nn.Module` with method `forward` taking `trigger` and `context` as arguments and returning a synthetic message :

```
.. code-block:: python
```

```
class CustomSynthesizer(torch.nn.Module):
    def forward(self, trigger, context):
        # synthesize the message
        return message
```

trigger will be a `torch.autograd.Variable` and context will be whatever is passed to the `dni.synthesizer_context` context manager, or `None` if `dni.synthesizer_context` is not used.

Example code for digit classification on MNIST using a CNN is at `examples/mnist-cnn` `<examples/mnist-cnn> _`.

Synthetic Gradient for Recurrent Networks ^^^

In this case we want to use DNI to approximate gradient from an infinitely-unrolled recurrent neural network and feed it to the last step of the RNN unrolled by truncated BPTT.

```
.. image:: images/rnn-update-unlock.png
```

We can use methods `make_trigger` and `backward` of `BackwardInterface` to do that:

```
.. code-block:: python
```

```
class Network(torch.nn.Module):

    def __init__(self):
        # ...

        # 1. create a BackwardInterface, assuming that dimensionality of
        #     the RNN hidden state is hidden_dim
        self.backward_interface = dni.BackwardInterface(
            dni.BasicSynthesizer(output_dim=hidden_dim, n_hidden=1)
        )

        # ...

    def forward(self, input, hidden):
        # ...

        # 2. call make_trigger on the first state of the unrolled RNN
        hidden = self.backward_interface.make_trigger(hidden)
        # run the RNN
        (output, hidden) = self.rnn(input, hidden)
        # 3. call backward on the last state of the unrolled RNN
        self.backward_interface.backward(hidden)
```

```

# ...

# in the training loop:
with dni.defer_backward():
    (output, hidden) = model(input, hidden)
    loss = criterion(output, target)
    dni.backward(loss)

```

`BackwardInterface.make_trigger` marks the first hidden state as a `trigger` used to update the gradient estimate. During the backward pass, gradient passing through the `trigger` will be compared to synthetic gradient generated based on the same `trigger` and the `Synthesizer` will be updated. `BackwardInterface.backward` computes synthetic gradient based on the last hidden state and backpropagates it.

Because we are passing both real and synthetic gradients through the same nodes in the computation graph, we need to use `dni.defer_backward` and `dni.backward`. `dni.defer_backward` is a context manager that accumulates all gradients passed to `dni.backward` (including those generated by `Interfaces`) and backpropagates them all at once in the end. If we don't do that, PyTorch will complain about backpropagating twice through the same computation graph.

Example code for word-level language modeling on Penn Treebank is at `examples/rnn` <`examples/rnn`> _.

Distributed training with a Complete Unlock ^^^

The paper describes distributed training of complex neural architectures as one of the potential uses of DNI. In this case we have a network split into parts A and B trained independently, perhaps on different machines, communicating via DNI. We can use methods `send` and `receive` of `BidirectionalInterface` to do that:

.. code-block:: python

```

class PartA(torch.nn.Module):

    def forward(self, input):
        x = input

        # ...

        # send the intermediate results computed by part A via DNI
        self.bidirectional_interface.send(x, input)

class PartB(torch.nn.Module):

    def forward(self, input):
        # receive the intermediate results computed by part A via DNI
        x = self.bidirectional_interface.receive(input)

        # ...

        return x

```

`PartA` and `PartB` have their own copies of the `BidirectionalInterface`. `BidirectionalInterface.send` will compute synthetic gradient w.r.t. `x` (intermediate results computed by `PartA`) based on `x` and `input` (input of the whole network), backpropagate it and update the estimate of `x`. `BidirectionalInterface.receive` will compute synthetic `x` based on `input` and in the backward pass, update the estimate of the gradient w.r.t. `x`. This should work as long as `BidirectionalInterface` parameters are synchronized between `PartA` and `PartB` once in a while.

There is no example code for this use case yet. Contributions welcome!